

EXP-7 Integrate a Vector Database with an LLM to build a Retrieval-Augmented Generation (RAG) system that answers questions based on external PDF documents.

AIM:

To build a Retrieval Augmented Generation (RAG) system that integrates a vector database with a language model to answer queries using external PDF documents.

DATASET:

Linkedin Profile PDF

PROCEDURE:

Step 1: Import Libraries

```
from langchain.document_loaders import PyPDFLoader
from langchain.text_splitter import RecursiveCharacterTextSplitter
from langchain.vectorstores import FAISS
from langchain.embeddings import OpenAIEmbeddings
from langchain.chains import RetrievalQA
from langchain.llms import OpenAI
```

Step 2: Load PDF

```
from google.colab import files
uploaded = files.upload()
loader = PyPDFLoader("your_file.pdf")
documents = loader.load()
```

Step 3: Text Chunking

```
text_splitter = RecursiveCharacterTextSplitter(
    chunk_size=500,
    chunk_overlap=50
)

texts = text_splitter.split_documents(documents)
```


Step 4: Create Embeddings

```
import os
os.environ["OPENAI_API_KEY"] = "your_api_key_here"

embeddings = OpenAIEmbeddings()
```

Step 5: Store in Vector Database (FAISS)

```
vector_db = FAISS.from_documents(texts, embeddings)
```

Step 6: Build Retrieval System

```
retriever = vector_db.as_retriever()
```

Step 7: Create RAG Chain

```
qa_chain = RetrievalQA.from_chain_type(
    llm=OpenAI(),
    retriever=retriever,
    return_source_documents=True
)
```

Step 8: Ask Questions

```
query = "What are the key skills mentioned in the document?"

result = qa_chain(query)

print("Answer:\n", result["result"])
```

OUTPUT:

The document highlights skills such as Python, Machine Learning, Data Analysis, and Deep Learning.

RESULT:

The RAG system successfully retrieved relevant content from the PDF and generated accurate, context aware answers to user queries.

