EXP 01 - Develop a predictive text system using N-Gram and HMM (Hidden Markov Models) to predict the next word based on previous context.

Aim:

To develop a predictive text system using N-Gram and HMM (Hidden Markov Models) to predict

the next word based on previous context.

Dataset: Brown Corpus

Step 1: Load the dataset

```
import nltk

nltk.download('brown')

from nltk.corpus import brown

sentences = brown.sents()
```

Step 2: Convert text to lowercase (remove special characters, punctuation, extra spaces)

```
import re

clean_sentences = []

for sent in sentences:
    text = " ".join(sent).lower()
    text = re.sub(r'[^a-z\s]', '', text)
    text = re.sub(r'\s+', ' ', text).strip()
    clean_sentences.append(text)
```

Step 3: Tokenization (sentence & word)

```
tokenized_sentences = []

for sent in clean_sentences:
    tokenized_sentences.append(sent.split())
```

Step 4: Create Vocabulary

```
vocab = set()

for sent in tokenized_sentences:
    for word in sent:
```

```
        vocab.add(word)
vocab_size = len(vocab)
```

N-GRAM MODEL

Step 5: Define N-gram (Bigram & Trigram):

```
from nltk.util import ngrams
bigrams = []
trigrams = []
for sent in tokenized_sentences:
    if len(sent) >= 2:
        bigrams.extend(list(ngrams(sent, 2)))
    if len(sent) >= 3:
        trigrams.extend(list(ngrams(sent, 3)))
```

Step 6: Generate sentence using N-gram:

```
generated_sentence = ['the']
for i in range(10):
    last_word = generated_sentence[-1]
    candidates = []
    for bg in bigrams:
        if bg[0] == last_word:
            candidates.append(bg[1])
    if len(candidates) == 0:
        break
    generated_sentence.append(candidates[0])
generated_sentence
```

o/p:

```
['the',
 'fulton',
 'county',
 'grand',
 'jury',
```

```
'said',
'friday',
'an',
'investigation',
'of',
'atlantas']
```

Step 7: Count word frequency:

```
from collections import Counter
unigram_freq = Counter()
bigram_freq = Counter()
for sent in tokenized_sentences:
    unigram_freq.update(sent)
    bigram_freq.update(list(ngrams(sent, 2)))
```

Step 8: Compute N-gram probability (MLE)

```
word1 = 'the'
word2 = 'man'
mle_prob = bigram_freq[(word1, word2)] / unigram_freq[word1]
mle_prob
```

o/p:

```
0.00265824412971088
```

Step 9: Apply Laplace smoothing

```
laplace_prob = (bigram_freq[(word1, word2)] + 1) / (unigram_freq[word1] + vocab_size)
laplace_prob
```

o/p:

```
0.0016086575021936238
```

Step 10: Predict next word

```
input_word = 'the'
max_prob = 0
```

```
next_word = None
for v in vocab:
    prob = (bigram_freq[(input_word, v)] + 1) / (unigram_freq[input_word] + vocab_size)
    if prob > max_prob:
        max_prob = prob
        next_word = v
next_word
```

o/p:

'first'

Step 11: Generate output

```
ngram_output = ['the', next_word]
ngram_output
```

o/p:

```
['the', 'first']
```

Accuracy for N-gram next-word prediction:

```
correct = 0
total = 0
for sent in tokenized_sentences[:1000]:
    if len(sent) < 2:
        continue
    prev_word = sent[-2]
    actual_word = sent[-1]
    max_prob = 0
    predicted_word = None
    for v in vocab:
        prob = (bigram_freq[(prev_word, v)] + 1) / (unigram_freq[prev_word] + vocab_size)
        if prob > max_prob:
            max_prob = prob
```

```
            predicted_word = v
    if predicted_word == actual_word:
        correct += 1
    total += 1
ngram_accuracy = correct / total
ngram_accuracy
```

o/p:

```
0.10040567951318459
```

HIDDEN MARKOV MODEL:

Step 5: Define HMM components

```
observations = vocab
transition = {}
emission = {}
states = set()
```

Step 6: POS Tagging (Hidden States)

```
nltk.download('averaged_perceptron_tagger')
tagged_sentences = []
for sent in tokenized_sentences[:5000]:
    tagged_sentences.append(nltk.pos_tag(sent))
```

Step 7: Estimate Transition & Emission Counts

```
from collections import Counter
transition_count = Counter()
emission_count = Counter()
state_count = Counter()
for sent in tagged_sentences:
    prev_tag = None
    for word, tag in sent:
        states.add(tag)
```

```
        state_count[tag] += 1
        emission_count[(tag, word)] += 1
        if prev_tag is not None:
            transition_count[(prev_tag, tag)] += 1
        prev_tag = tag
```

Step 8: Calculate Transition & Emission Probabilities

```
for (t1, t2) in transition_count:
    transition[(t1, t2)] = transition_count[(t1, t2)] / state_count[t1]
for (tag, word) in emission_count:
    emission[(tag, word)] = emission_count[(tag, word)] / state_count[tag]
```

Step 9: Train HMM Model

```
model_states = list(states)
```

Step 10: Predict Next Word

```
input_sentence = ['the', 'man']
last_word = input_sentence[-1]
candidate_words = {}
for (tag, word) in emission:
    if word == last_word:
        for next_tag in model_states:
            if (tag, next_tag) in transition:
                for w in observations:
                    if (next_tag, w) in emission:
                        candidate_words[w] = transition[(tag, next_tag)] * emission[(next_tag, w)]
hmm_next_word = max(candidate_words, key=candidate_words.get)
hmm_next_word
```

Step 11: Generate Output

```
hmm_output = input_sentence + [hmm_next_word]
hmm_output
```

Accuracy for HMM next-word prediction (Emission-based)

```
correct = 0
total = 0


for sent in tagged_sentences[:500]:
    if len(sent) < 2:
        continue
    prev_word, prev_tag = sent[-2]
    actual_word, actual_tag = sent[-1]
    predictions = {}
    for w in vocab:
        if (prev_tag, w) in emission:
            predictions[w] = emission[(prev_tag, w)]
    if len(predictions) == 0:
        continue
    predicted_word = max(predictions, key=predictions.get)
    if predicted_word == actual_word:
        correct += 1
    total += 1
hmm_accuracy = correct / total
hmm_accuracy
```

O/P:

0.0

Result:

N-gram model achieved about **10% accuracy**, as it predicts the next word based on frequent word sequences in the corpus. HMM model showed **0% accuracy**, since it predicts words based on POS patterns, which rarely match the exact next word in the dataset.